


Leader Election



```
leader_election.py — leader_election
leader_election.py
1 import etcd3
2
3 import sys
4 import time
5 from threading import Event
6
7
8 # The current leader is going to be the value with this key.
9 LEADER_KEY = "/algoexpert/leader"
10
11
12 # Entrypoint of the program.
13 def main(server_name):
14     # Create a new client to etcd.
15     client = etcd3.client(host="localhost", port=2379)
16
17     while True:
18         is_leader, lease = leader_election(client, server_name)
19
20         if is_leader:
21             print("I am the leader.")
22             on_leadership_gained(lease)
23         else:
24             print("I am a follower.")
25             wait_for_next_election(client)
26
27
28 # This election mechanism consists of all clients trying to put their name
29 # into a single key, but in a way that only works if the key does not
30 # exist (or has expired before).
31 def leader_election(client, server_name):
32     print("New leader election happening.")
33     # Create a lease before creating a key. This way, if this client ever
34     # lets the lease expire, the keys associated with that lease will all
35     # expire as well.
36     # Here, if the client fails to renew lease for 5 seconds (network
37     # partition or machine goes down), then the leader election key will
38     # expire.
39     # https://help.compose.com/docs/etcd-using-etcd3-features#section-leases
40     lease = client.lease(5)
41
42     # Try to create the key with your name as the value. If it fails, then
43     # another server got there first.
```

```
leader_election — bash — 0.0 x 24
Clements-MBP:leader_election clementmihailescu$ python3 leader_election.py server2

leader_election — bash — 0.0 x 24
Clements-MBP:leader_election clementmihailescu$ python3 leader_election.py server4
```

This code example was based off
of the code found at:
<https://www.sandtable.com/etcd3-leader-election-using-python/>

The image displays a development environment with a code editor and a terminal window.

Code Editor: The file `leader_election.py` contains the following Python code:

```
59 except Exception:
60     # Here we most likely got a client timeout (from
61     # network issue). Try to revoke the current lease
62     # so another member can get leadership.
63     lease.revoke()
64     return
65 except KeyboardInterrupt:
66     print("\nRevoking lease; no longer the leader.")
67     # Here we're killing the process. Revoke the lease and exit.
68     lease.revoke()
69     sys.exit(1)
70
71
72 def wait_for_next_election(client):
73     election_event = Event()
74
75     def watch_callback(resp):
76         for event in resp.events:
77             # For each event in the watch event, if the event is a deletion
78             # it means the key expired / got deleted, which means the
79             # leadership is up for grabs.
80             if isinstance(event, etcd3.events.DeleteEvent):
81                 print("LEADERSHIP CHANGE REQUIRED")
82                 election_event.set()
83
84     watch_id = client.add_watch_callback(LEADER_KEY, watch_callback)
85
86     # While we haven't seen that leadership needs change, just sleep.
87     try:
88         while not election_event.is_set():
89             time.sleep(1)
90     except KeyboardInterrupt:
91         client.cancel_watch(watch_id)
92         sys.exit(1)
93
94     # Cancel the watch; we see that election should happen again.
95     client.cancel_watch(watch_id)
96
97
98 # Try to insert a key into etcd with a value and a lease. If the lease expires
99 # that key will get automatically deleted behind the scenes. If that key
100 # was already present, this will raise an exception.
```

Terminal Window: The terminal shows the command to run the script and the output:

```
~/Documents/Content/Design_Fundamentals/Examples/leader_election -- -bash
Clements-MBP:leader_election clementmihailescu$ python3 leader_election.py server2
```

The terminal window is titled `leader_election -- -bash` and shows the command `python3 leader_election.py server2` being executed. The output shows the script running successfully.

The image displays four terminal windows arranged in a 2x2 grid, showing the execution of a distributed system simulation. Each window has a title bar indicating the current directory and the command being run.

- Top-left window:** Title is `leader_election - Python 3 leader.py server1 -- 93x24`. The output shows a server refreshing leases and then revoking them. The command is `python3 leader_election.py server1`.
- Top-right window:** Title is `leader_election - Python 3 leader.py server2 -- 93x24`. The output shows a server starting a new leader election. The command is `python3 leader_election.py server2`.
- Bottom-left window:** Title is `leader_election - Python 3 leader.py server3 -- 93x24`. The output shows a server becoming a leader after a leadership change. The command is `python3 leader_election.py server3`.
- Bottom-right window:** Title is `leader_election -- -bash -- 93x24`. The output shows a server refreshing leases continuously. The command is `-bash`.

5 Prerequisites

Availability

The odds of a particular server or service being up and running at any point in time, usually measured in percentages. A server that has 99% availability will be operational 99% of the time (this would be described as having two **nines** of availability).

High Availability

Used to describe systems that have particularly high levels of availability, typically 5 nines or more; sometimes abbreviated "HA".

Redundancy

The process of replicating parts of a system in an effort to make it more reliable.

Strong Consistency

Strong Consistency usually refers to the consistency of ACID transactions, as opposed to **Eventual Consistency**.

Eventual Consistency

A consistency model which is unlike **Strong Consistency**. In this model, reads might return a view of the system that is stale. An eventually consistent datastore will give guarantees that the state of the database will eventually reflect writes within a time period (could be 10 seconds, or minutes).

5 Key Terms

Leader Election

The process by which nodes in a cluster (for instance, servers in a set of servers) elect a so-called "leader" amongst them, responsible for the primary operations of the service that these nodes support. When correctly implemented, leader election guarantees that all nodes in the cluster know which one is the leader at any given time and can elect a new leader if the leader dies for whatever reason.

Consensus Algorithm

A type of complex algorithms used to have multiple entities agree on a single data value, like who the "leader" is amongst a group of machines. Two popular consensus algorithms are **Paxos** and **Raft**.

Paxos & Raft

Two consensus algorithms that, when implemented correctly, allow for the synchronization of certain operations, even in a distributed setting.

Etcd ⚡

Etcd is a strongly consistent and highly available key-value store that's often used to implement leader election in a system.

ZooKeeper ⚡

ZooKeeper is a strongly consistent, highly available key-value store. It's often used to store important configuration or to perform leader election.